




# OCTOPUS

GROW TENTACLES AND CONTROL MACHINES



# OCTOPUS

Configuring a machine (virtual or not) is easy, you can do what you want manually.

Configuring a few machines can be tedious to do manually.

Configuring manually a lot of machines is exhausting, and not worth of investing that much time into it.

There comes the solution, for us perfectly lazy DevOps engineers (in the making): Ansible!

Ansible is an automation tool that allows one to automate an other part of the duty of operators and developers: *configuration management*.

Configuration management refers to the tools and practices used to manage the state in which machines needs to be, in a given set of machines.

Ansible is an incredibly handy tool that allows you to do that in an efficient and scalable way.



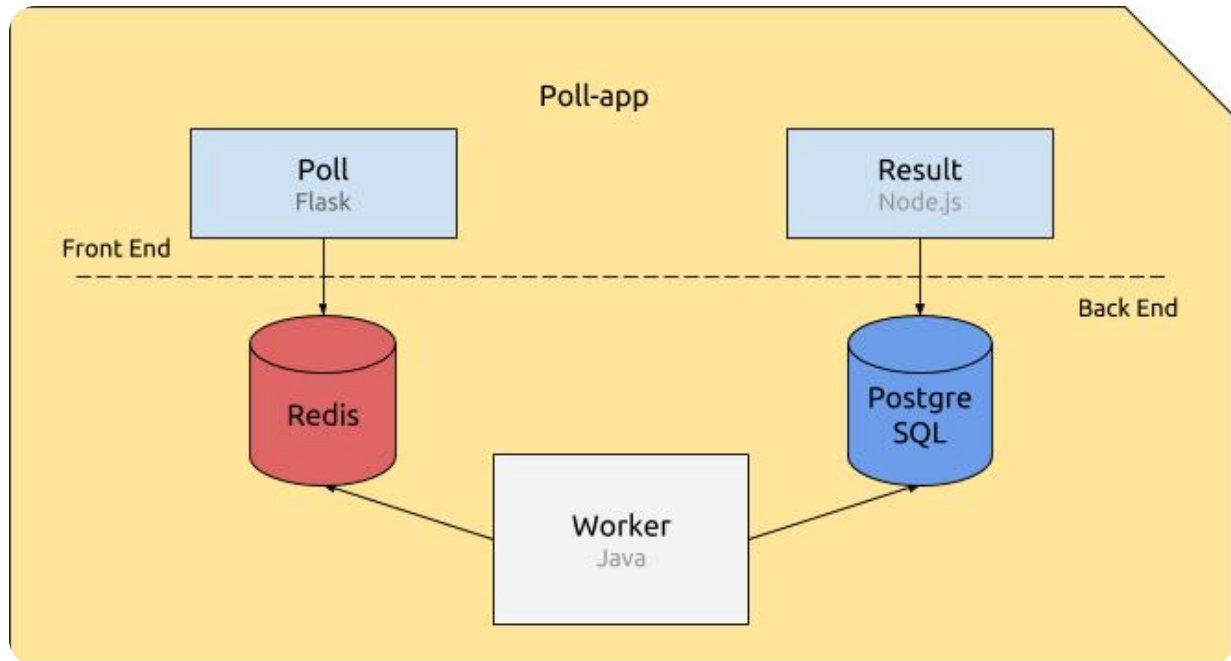
# General objective

You are going to work again with the poll application you worked with during the Popeye project.

As a reminder, the poll application starts with a Python Flask web client that gathers the votes, and then pushes them into a Redis queue.

Afterwards, the Java worker consumes the votes stored in the Redis queue, and pushes them into a PostgreSQL database.

Finally, the Node.js web client fetches the votes from the PostgreSQL database and displays the results.



This time, you are going to *deploy* the application onto 5 different machines, without using containers, by using Ansible.

# Technical details

For this project, you have to turn in a `playbook.yml` file and a directory named `roles` that will hold your Ansible roles (see section *Repository structure* below).

Your project will be manually evaluated.

It will be tested using the following commands:

```
Terminal
~/B-DOP-400> export ANSIBLE_VAULT_PASSWORD_FILE=/tmp/.vault_pass
B-DOP-400> echo verySecretPassword > /tmp/.vault_pass
B-DOP-400> ansible-playbook -i production playbook.yml
```



We will use our own inventory file named `production`, which will ensure that your playbook can access the machines on which we will test your project. It will be in compliance with the requirements described in the *Environment* section below.

**The inventory file will be configured to connect to the machines as a normal user (i.e.: not `root`) which itself will be able to use the `sudo` command.**



Docker, Podman, other containerization technologies, and Ansible Galaxy are strictly **forbidden in all their forms** for this project.

You may however use the remaining [Community namespace's collections](#).



**Security cannot be ignored when working in a DevOps environment. As such, any clear-text password found in your repository will cause the evaluation to stop right away, and the entire project will be considered as failed. You have been warned.**

# Features

## Roles

You have to write 6 roles.

`base`

---

This role is associated to all machines.

- ✓ Installs useful packages (such as Git) using `apt`.
- ✓ Configures instance.



Think about what is necessary to do to ensure a properly configured instance.  
*What about useful **commands** or **system tools**? ;)*



Do not install useless or too big packages.

`redis`

---

- ✓ Installs Redis.
- ✓ Sets up Redis.

## postgresql

---

- ✓ Installs PostgreSQL 16.
- ✓ Creates a user `paul` with the password `democracyIsFragile` and limited permissions.
- ✓ Creates the schema of the database `paul`.



Be careful!

The `paul` **user** and the `paul` **database** are two different things!

The `paul` user is a user created in the database, not on the operating system itself.

## poll

---

- ✓ Uploads `poll` service.
- ✓ Installs dependencies.
- ✓ Runs the poll web client.

## worker

---

- ✓ Uploads `worker` service.
- ✓ Installs dependencies.
- ✓ Builds the worker.
- ✓ Runs the worker.

## result

---

- ✓ Uploads `result` service.
- ✓ Installs dependencies.
- ✓ Runs the result web client.



Avoid using the `command` / `shell` / `raw` modules as much as possible.  
Have a look at dedicated and cleaner modules instead, such as `apt_key`, `apt_repository`, or `pip`.  
Use the `command` module only if no dedicated module exists.

All services must be managed by `systemd` and start automatically on boot.

In order to respect the `12 factor` best practices, services must be configured with environment variables, such as: host, port, user, password, database name, etc.



After applying an Ansible playbook twice, you should have as few `changed` tasks as possible in your `PLAY RECAP` (see the example below).  
This is the famous notion of *idempotence* that Ansible is famous for, and which you should strive for.

```
PLAY RECAP *****
1.1.1.1      : ok=49  changed=0  unreachable=0  failed=0  skipped=12  rescued=0  ignored=0
2.2.2.2      : ok=25  changed=0  unreachable=0  failed=0  skipped=16  rescued=0  ignored=0
3.3.3.3      : ok=42  changed=1  unreachable=0  failed=0  skipped=22  rescued=0  ignored=0
```

## Environment

Your inventory **must** have the following 5 groups, each containing 1 instance:

- ✓ `redis` with a `redis-1` instance;
- ✓ `postgresql` with a `postgresql-1` instance;
- ✓ `poll` with a `poll-1` instance;
- ✓ `result` with a `result-1` instance;
- ✓ and `worker` with a `worker-1` instance.

You will need 5 virtual machines based on Debian 12.

You can run it locally, it is however recommended to use a cloud platform.

Do not spend too much credit however, you may also need cloud platforms for the next DevOps projects.

## Repository structure

In the end, your repository must at least contain the following files:

```
|-- playbook.yml
|-- poll.tar
|-- result.tar
|-- worker.tar
|-- group_vars
|  |-- all.yml
|-- roles
|   |-- base
|   |  |-- tasks
|   |  |  |-- main.yml
|   |-- postgresql
|   |  |-- files
|   |  |  |-- pg_hba.conf
|   |  |  |-- schema.sql
|   |  |-- tasks
|   |  |  |-- main.yml
|   |-- redis
|   |  |-- files
|   |  |  |-- redis.conf
|   |  |-- tasks
|   |  |  |-- main.yml
|   |-- poll
|   |  |-- files
|   |  |  |-- poll.service
|   |  |-- tasks
|   |  |  |-- main.yml
|   |-- result
|   |  |-- files
|   |  |  |-- result.service
|   |  |-- tasks
|   |  |  |-- main.yml
|   |-- worker
|   |  |-- files
|   |  |  |-- worker.service
|   |  |-- tasks
|   |  |  |-- main.yml
```



# Bonus

You want to go further? Great! Here is an idea:

- ✓ make Redis more secure by adding a password to it (you are allowed to update the applications' code in this case).

{EPITECH}

